

Design Decisions:

1. I only scraped English reviews with the region in the US.
2. I prioritized newest reviews, which mimics the real scraping process, but it lacks certain more helpful reviews, so I think we can combine the current newest reviews with helpful reviews for more usefulness.
3. Removed three columns (replyContent, repliedAt, appVersion) that contained a great amount of null values, but it is not quite helpful for our later sentiment analysis, so I dropped them for clear dataset features.
4. Added an app identifier in the first column of the data to identify each app.
5. Defined low-quality reviews via three metrics: short character length, pure-emoji-only content, and entries with no alphabetic words, and also quantifies the proportion of non-actionable feedback to evaluate dataset efficiency.

Limitations:

1. Sampling Bias
The scraping only uses "sort.newest" for primary data collection. We did not evenly sample reviews sorted by helpfulness, which creates bias toward recently posted feedback and gets some less meaningful reviews.
2. Repeated Reviews
We have 5002 repeated reviews, which takes 1/3 of the entire dataset, and I think that is also relevant to the sampling bias mentioned above.
3. Meaningless Reviews
We have in total 9918 short reviews, which may include reviews that only has emojis, and reviews that have no actual word in it. This takes up to 1/2 of the entire dataset, which is a concern because we want more information about products and add more commercial value to it.